

Decentralized & Collaborative AI on Blockchain

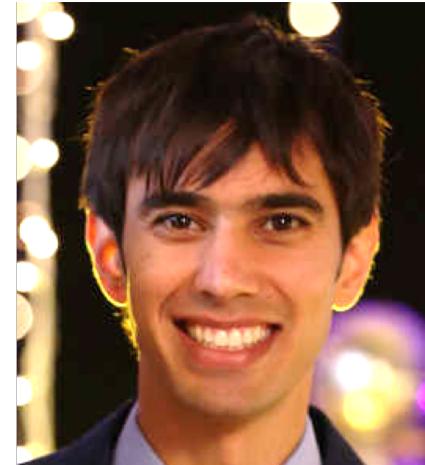
Crowdsourcing and Machine Learning
Models on the Blockchain

Justin Harris
Bo Waggoner

TEAM



Justin D. Harris
justin.harris@microsoft.com
Senior Software Engineer
Microsoft Research, Montreal



Bo Waggoner (bwag@colorado.edu)
Former Post-doc, Microsoft Research, NYC
Assistant professor, University of Colorado,
Boulder.

OUTLINE

Explanation

Why Blockchain

ML Models

Economics

Demo

THE PROBLEMS

- Difficult to set up AI/ML systems
 - hardware constraints → cost constraints
 - scaling
- Difficult to keep an up-to-date model deployed
- Advancements are centralized
 - datasets are not shared
 - charged per-query

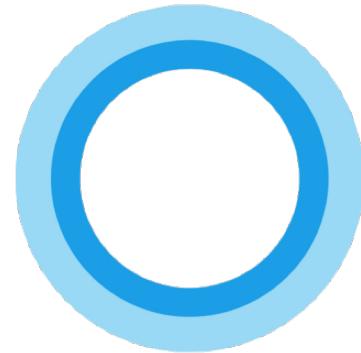
THE SOLUTION



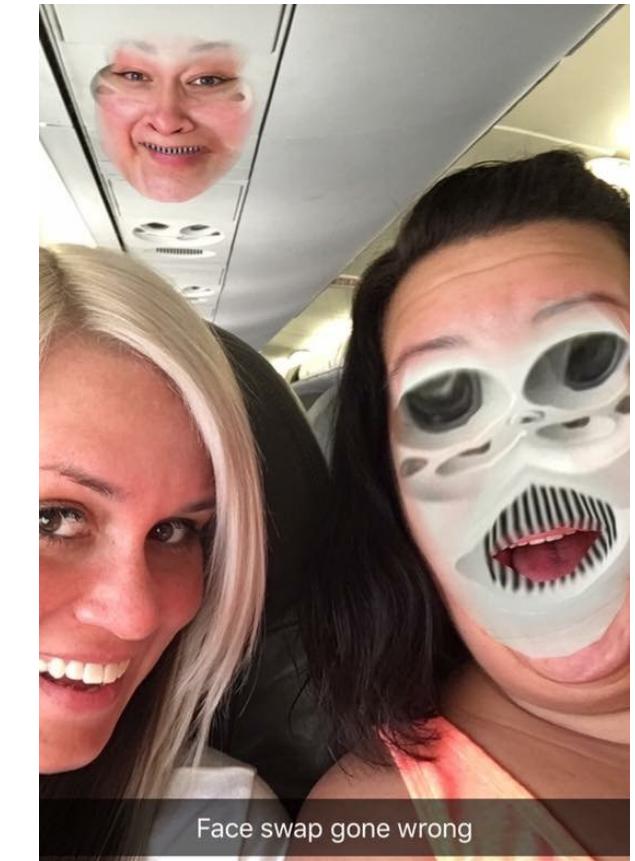
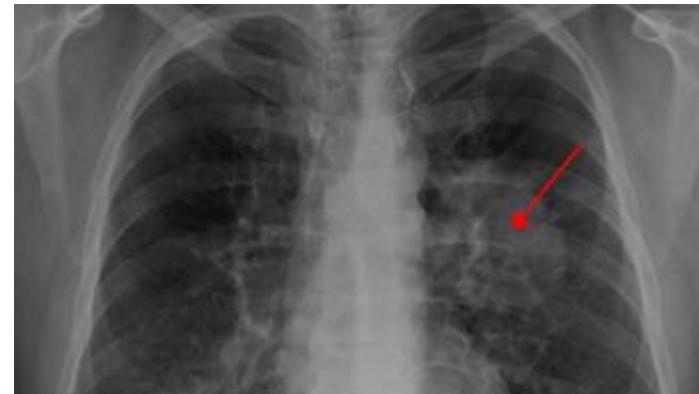
Store models on a blockchain.

MOTIVATIONS

- Improve models you use as data evolves
- Crowdsourcing: access people + data



Hi. I'm Cortana.

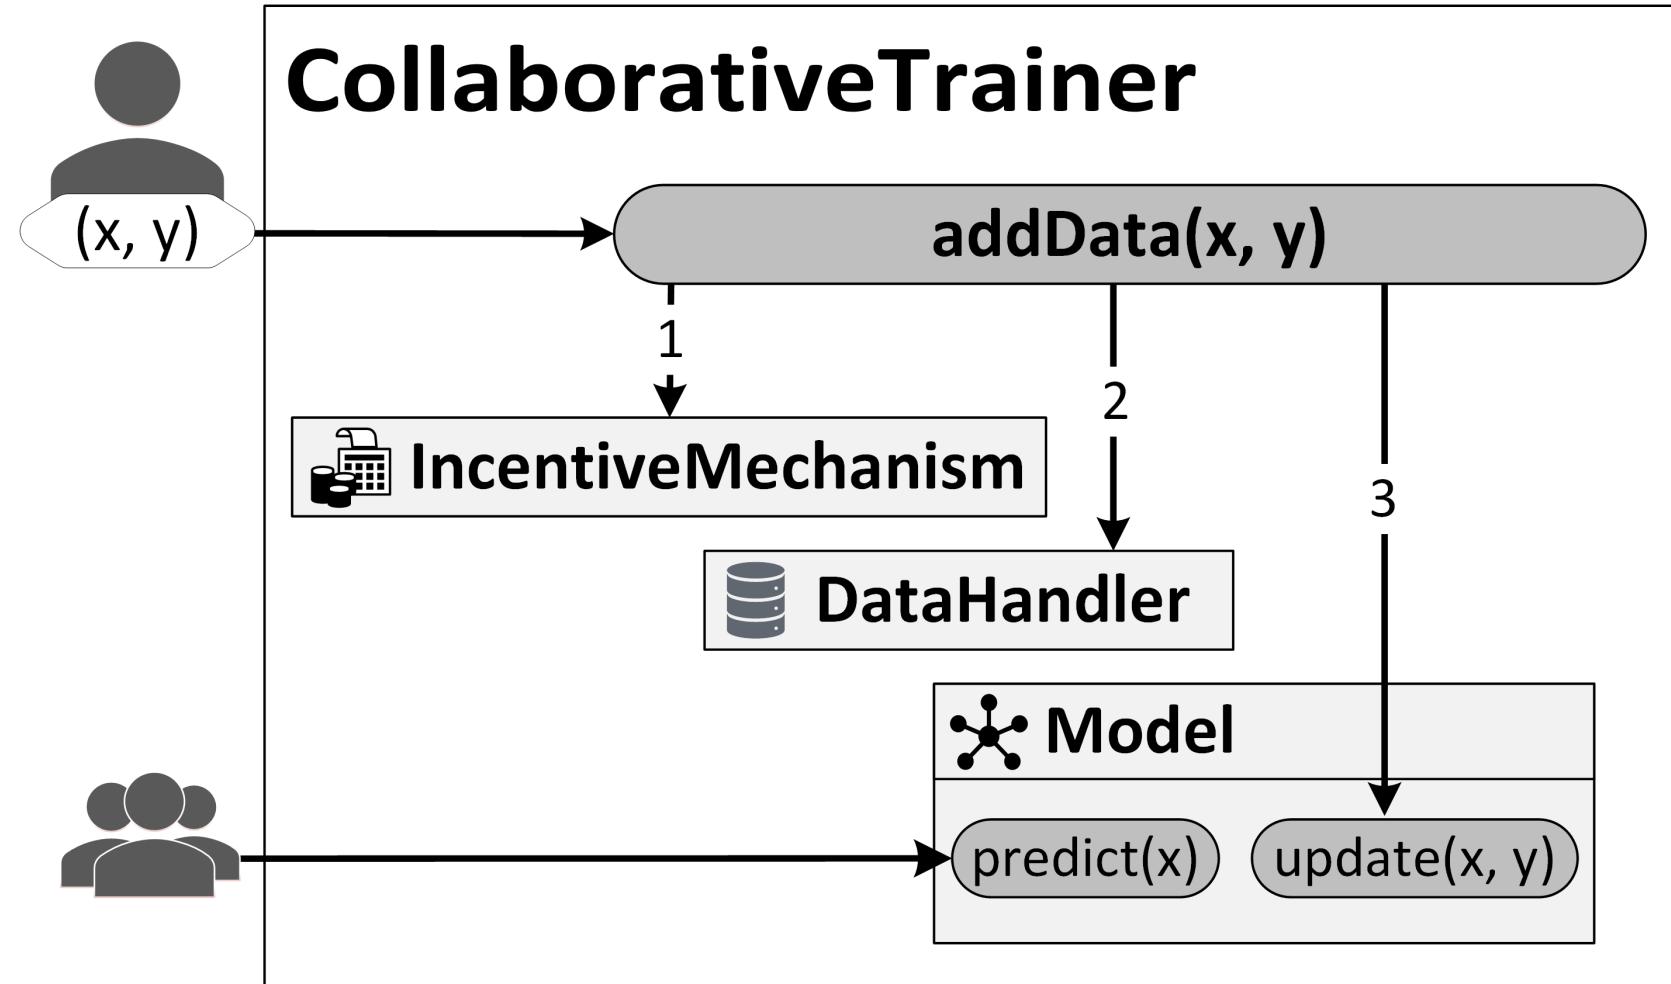


Face swap gone wrong

OVERVIEW

Adding data consists
of 3 steps.

- Boxes are smart contracts
- Rounded boxes are methods



BLOCKCHAIN

- Public, Persistent, and Decentralized
- Versioned
- Models evolve over time
- Transparency & Trust
- Simplified payment

MODELS

- Supervised machine learning: data with labels
- Minimize gas costs → efficient to train models
 - E.g. Perceptron, Nearest Centroid Classifier, or Naïve Bayes
- Encoding “off-chain” then fine-tuning “on-chain”
 - E.g. Image recognition

Ethereum Gas Costs

Action	Gas Cost	USD ¹
Deploy model contract of Perceptron with 100 weights	3,845,840	\$4.71
Add data with 15 words (model agrees) ²	177,693	\$0.22
Add data with 15 words (model disagrees) ²	249,037	\$0.30

¹Approximate costs in July 2019 with a modest gas price of 4gwei.

²Perceptron models are only updated when the model disagrees.

MODELS

- Perceptron:
 - Only update when expected class \neq predicted class
 - Easy to update: $w(t + 1) = w(t) + r \cdot (y - \hat{y}) \cdot x$
 - $w(t)$: weights at time t
 - r : learning rate
 - $\hat{y} = w(t) \cdot x + b$: current classification
 - y : expected classification
- Nearest Centroid Classifier
 - Easy to update moving average: $\text{avg}(t + 1) = \frac{x + n \cdot \text{avg}(t)}{t + 1}$
 - Enforce normalized: no one can move a centroid too much
 - Encode “off-chain” using a known encoder: tested with 512 dimensions

INCENTIVIZING QUALITY DATA

There are many ways to encourage contributors to submit good quality data.
We analyze several examples in our paper:

1. **Gamified** (non-financial, points + badges like Stackoverflow)
2. Based on established theory in **Prediction Markets**
3. Deposit, Refund, and Take: **Self-Assessment** (screenshot demo later)

INCENTIVIZING QUALITY DATA GAMIFIED (NON-FINANCIAL)

- Points + badges like Stackoverflow
 - Milestones for number of contributions
 - Points for submitting diverse data
 - Badges for using different labels
 - Extra points for submitting data frequently
- Can be tracked on-chain or off-chain by 3rd parties



INCENTIVIZING QUALITY DATA BASED ON PREDICTION MARKETS

Prediction Market: Bet or contribute a belief on the outcome of an event.
E.g. Winner of a soccer game or an election.

Here we use ideas from prediction markets to incentivize good data contributions such as in "[A Collaborative Mechanism for Crowdsourcing Prediction Problems](#)" (Abernethy et al., NeurIPS 2011) and "[A Market Framework for Eliciting Private Data.](#)" (Waggoner et al., NeurIPS 2015).

Phases

- 1) Commitment
- 2) Participation
- 3) Reward

INCENTIVIZING QUALITY DATA BASED ON PREDICTION MARKETS

1) Commitment Phase

- A generous **provider** stakes a **bounty** to be split and rewarded to contributors.
- Now the provider must prove they have a valid **test set** but without revealing all of it yet.¹
 - Provider shares hashes for portions of their test set: h_1, h_2, \dots, h_N
 - Provider reveals a portion of the test set randomly chosen by a smart contract:
 $H = \{h_i : 1 \leq i \leq N\}, |H| < N$

¹Similar to the DanKu Protocol: <https://algorithmia.com/research/ml-models-on-blockchain>

INCENTIVIZING QUALITY DATA BASED ON PREDICTION MARKETS

2) Participation Phase

- **Participants** submit one training **data** sample at time along with a small **deposit** of funds.
- The shared **model** is **updated** using the provided data sample.

INCENTIVIZING QUALITY DATA BASED ON PREDICTION MARKETS

3) Reward Phase

- The **provider reveals** the rest of the **test set** and the smart contract validates that it matches the hashes they originally gave in the Commitment Phase.
- Participants are **rewarded** based on how much their data contribution helped the model **improve** its **accuracy** on the test set:

change in loss (error rate): $L(h_{t-1}, D) - L(h_t, D)$

change in accuracy: $A(h_t, D) - A(h_{t-1}, D)$

INCENTIVIZING QUALITY DATA BASED ON PREDICTION MARKETS

B = bounty

Amount distributed $\leq B \cdot [L(h_0, D) - L(h_T, D)]$

```
Let  $b_t = 1$  for all  $t$       // balance initially equals stake
Let list  $S = (1, \dots, T)$  // list initially contains everyone
for  $i = 1, \dots, B$  do
    for each participant  $t$  in  $S$  do
        Let  $t'$  be previous participant in  $S$ , or 0 if none.
        Participant  $t$ 's balance is changed:
```

$$b_t \leftarrow b_t + L(h_{t'}, D) - L(h_t, D)$$

```
Let list  $S = (t \in S : b_t \geq 1)$ . // all who can re-stake
1 stay in  $S$ 
Each participant  $t$  is paid  $b_t$ .
```

BASED ON PREDICTION MARKETS

```
Let  $b_t = 1$  for all  $t$  // balance initially equals stake  
Let list  $S = (1, \dots, T)$  // list initially contains everyone  
for  $i = 1, \dots, B$  do
```

```
    for each participant  $t$  in  $S$  do
```

```
        Let  $t'$  be previous participant in  $S$ , or 0 if none.
```

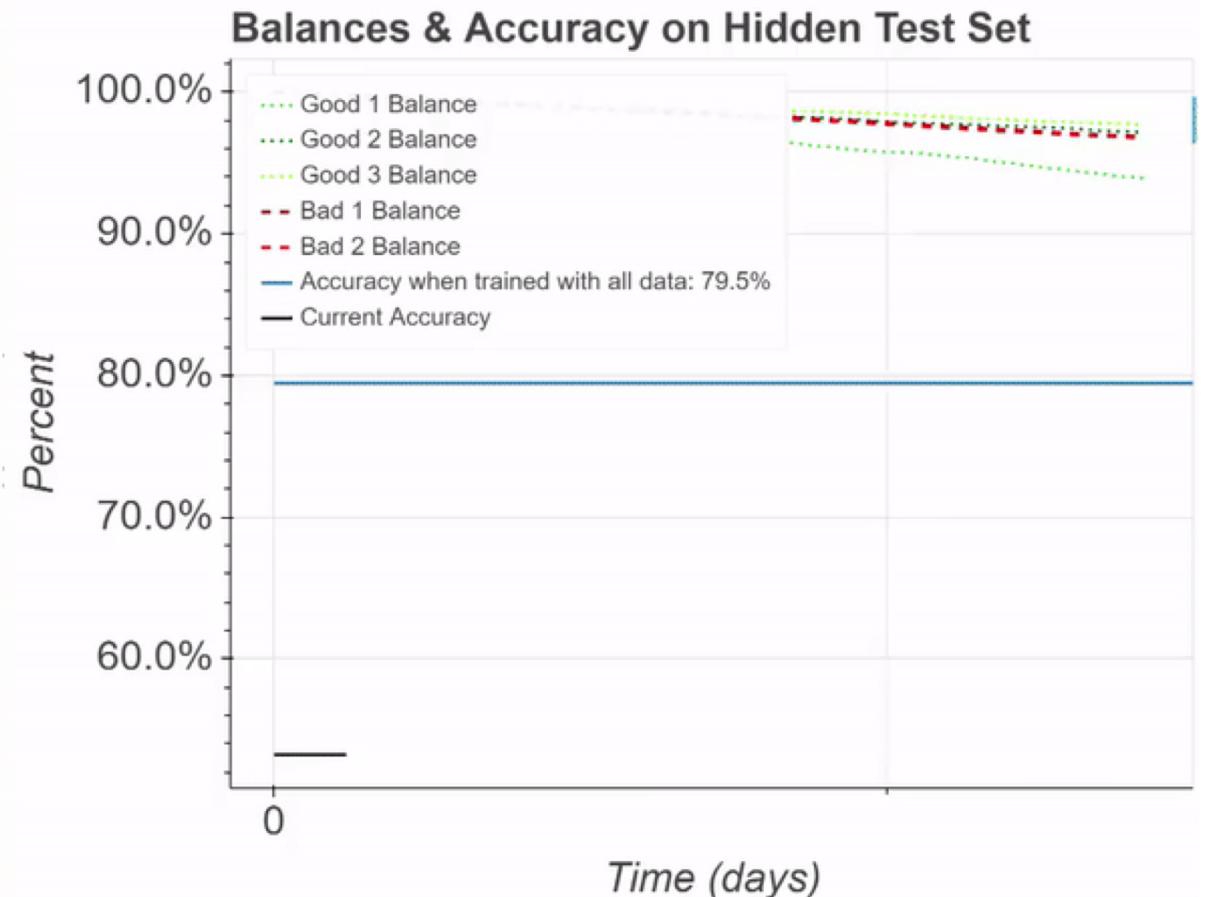
```
        Participant  $t$ 's balance is changed:
```

$$b_t \leftarrow b_t + L(h_{t'}, D) - L(h_t, D)$$

```
    Let list  $S = (t \in S : b_t \geq 1)$ . // all who can re-stake  
    1 stay in  $S$ 
```

```
    Each participant  $t$  is paid  $b_t$ .
```

Amount distributed $\leq B \cdot [L(h_0, D) - L(h_T, D)]$



Participants submit data and deposits.

INCENTIVIZING QUALITY DATA DEPOSIT, REFUND, AND TAKE: SELF-ASSESSMENT

Demo

- Predict
- Deposit
- Refund
- Take

DEMO

Democratize AI

+ [CREATE NEW MODEL](#) 

[IMDB Review Sentiment Model](#)
82.9% 

[Self-driving Car](#)
87.2%

[Cancer Fighting Nanobots](#)
75.7%

[Bengio AGI](#)
99.1%

DEMO

IMDB Review Sentiment Model

A simple IMDB sentiment analysis model

Refund Time: a few seconds

Claim Time: a few seconds

Current Required Deposit: £0.105882

PREDICT

TRAIN

REFUND

REWARD



Input

Great movie!

GET PREDICTION



Prediction: Positive Sentiment

DEMO

IMDB Review Sentiment Classifier

A simple IMDB sentiment analysis model.

Your score: 100.00% (1/1)
Refund Time: a few seconds
Claim Time: a few seconds
Current Required Deposit: ⓢ0.450000

PREDICT TRAIN REFUND REWARD

Data Sample
This was the best movie ever!

Classification
Positive

TRAIN

Deposit $\propto 1/\Delta t$



MetaMask Notification Local 7545

Good → 0x6417...14...

CONTRACT INTERACTION

0.45

DETAILS DATA EDIT

GAS FEE 0 No Conversion Rate Available

AMOUNT + GAS FEE

TOTAL 0.45 No Conversion Rate Available

Reject Confirm



DEMO

IMDB Review Sentiment Model

A simple IMDB sentiment analysis model



Your score: 100.00% (2/2)

Refund Time: a few seconds

Claim Time: a few seconds

Current Required Deposit: ₩0.036735

PREDICT

TRAIN

REFUND

REWARD

Data	Classification	Initial Deposit	Date Added	
"this was the best movie i've ever seen"	Positive Sentiment	₩0.087805	Wed Nov 21 2018 13:59:34 GMT-0500 (Eastern Standard Time)	Already refunded or completely claimed.
"great film"	Positive Sentiment	₩0.900000	Wed Nov 21 2018 14:00:24 GMT-0500 (Eastern Standard Time)	Already refunded or completely claimed.
"best movie ever"	Positive Sentiment	₩0.720000	Wed Nov 21 2018 14:00:05 GMT-0500 (Eastern Standard Time)	REFUND ₩0.720000



DEMO

IMDB Review Sentiment Model

A simple IMDB sentiment analysis model

Refund Time: a few seconds
Claim Time: a few seconds
Current Required Deposit: ₩0.036364



PREDICT	TRAIN	REFUND	RWARD
Data	Classification	Initial Deposit	Date Added
"bad movie"	Positive Sentiment	₩0.257143	Wed Nov 21 2018 14:21:23 GMT-0500 (Eastern Standard Time) Classification doesn't match. Got "Negative Sentiment".
"most amazing drama ever"	Negative Sentiment	₩0.112500	Wed Nov 21 2018 14:18:03 GMT-0500 (Eastern Standard Time) Already refunded or completely claimed.

DEMO

IMDB Review Sentiment Model

A simple IMDB sentiment analysis model

Your score: 100.00% (2/2)
Refund Time: a few seconds
Claim Time: a few seconds
Current Required Deposit: ₩0.036000



PREDICT	TRAIN	REFUND	REWARD
Data	Classification	Initial Deposit	Date Added
"most amazing drama ever"	Negative Sentiment	₩0.112500	Wed Nov 21 2018 14:18:03 GMT-0500 (Eastern Standard Time)
"bad movie"	Positive Sentiment	₩0.257143	Wed Nov 21 2018 14:21:23 GMT-0500 (Eastern Standard Time)

TAKE ₩0.257143

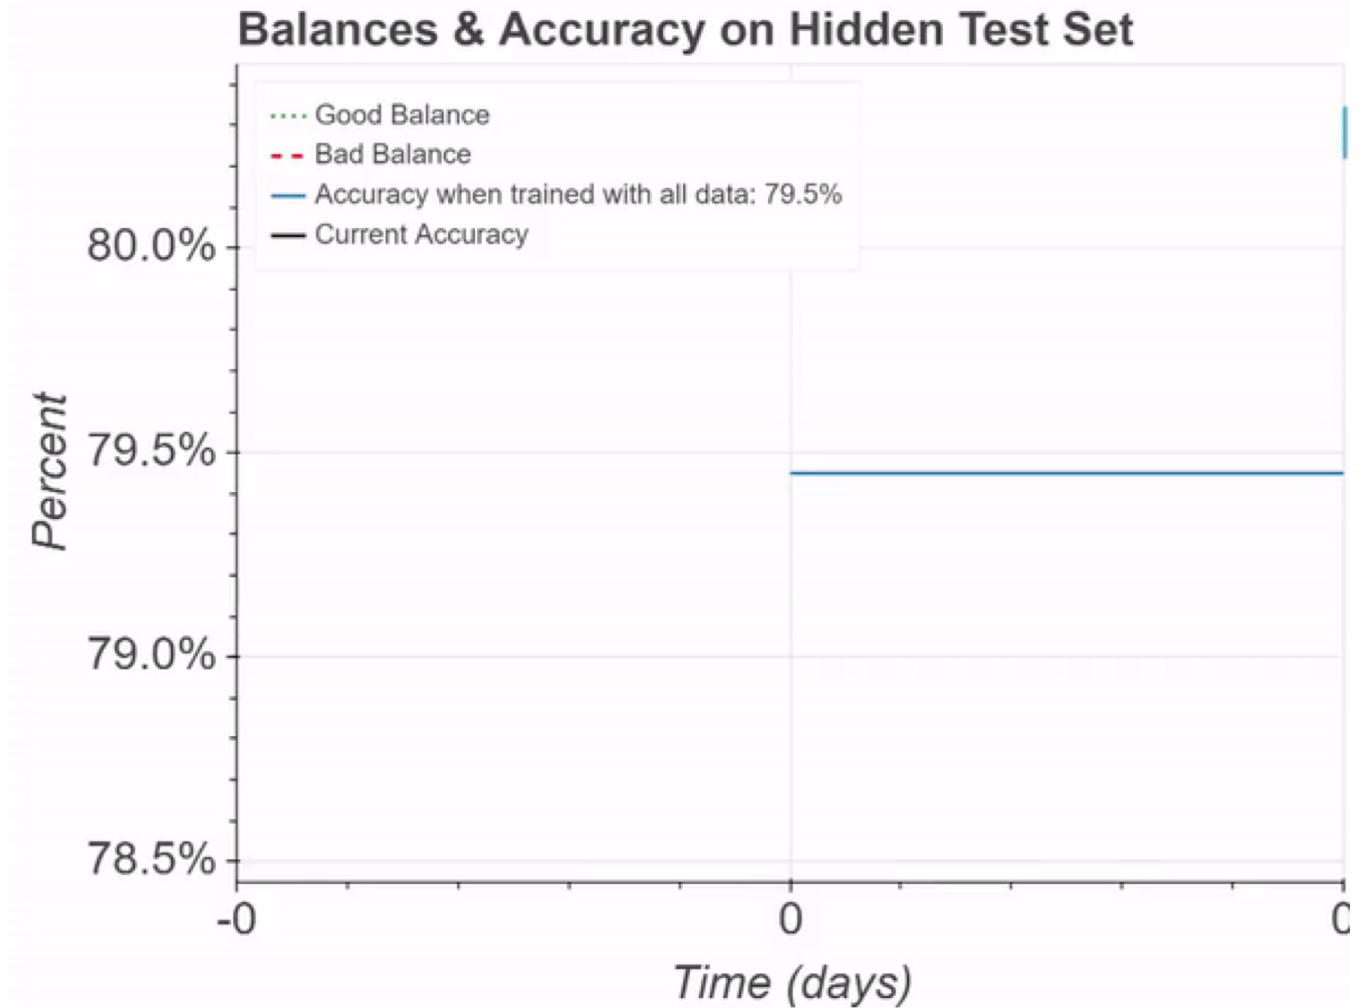


SIMULATION

“Bad Agent” frequently adds incorrect data.

The model can still maintain accuracy.

Honest contributors can still profit.



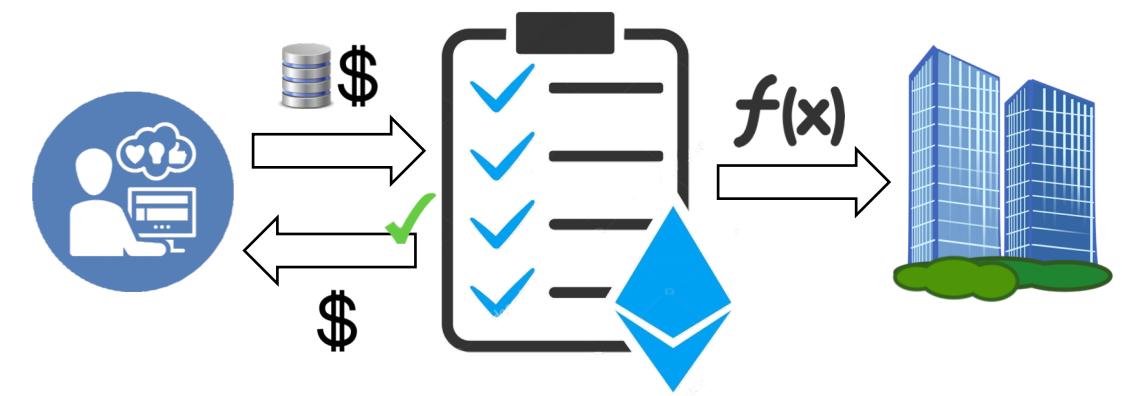
SUMMARY

Goals:

- Free to use models in smart contracts
- Build high quality datasets

Method:

- Deploy an initial model
- Contributors submit data + deposit
- Contributors can get a reward after submitting good data
- The model remains free to use for inference



NEXT STEPS

- Analyze more incentive mechanisms
- Find the best models to use
- Privacy: handle private data?
- Off-chain models?
- 3rd Party Platforms: free to contribute
- Unsupervised techniques for filtering bad data



blog: <https://aka.ms/0xDeCA10B-blog1>
<https://github.com/microsoft/0xDeCA10B>